

Aman Sharma

+91 9871987731  GitHub  LinkedIn  amands.5153@gmail.com

SUMMARY

Grounded in classical Data Science — statistics, forecasting, and machine learning — now building deep expertise in LLM fine-tuning, RAG pipelines, and agentic systems. Driven towards building production-grade AI systems.

EDUCATION

Indian Institute of Information Technology, Lucknow

Expected August 2027

M.Sc. Data Science | SGPA: 8.65/10 (Semester 1)

EXPERIENCE

AI Agent Developer Intern

May 2026 – Present

Genotek Global

Remote

- Built and deployed a **Telegram bot** on **AWS EC2** for field sales — GPS-verified check-in/check-out with a 3-gate anti-cheat system (live location · freshness · accuracy · Haversine distance), Google Places verification, **Supabase** backend, and HubSpot CRM sync; validated end-to-end with company leadership
- Shipped an LLM meeting-debrief pipeline — local **faster-whisper** (CPU int8, Hinglish-aware) → **Groq LLaMA 3.1 70B** → structured HubSpot engagements — atop a **15-state python-telegram-bot FSM** with multi-rep session isolation and drift classification (clean/suspicious/flagged) + manager alerting

Data Science Intern

Sep 2025 – Nov 2025

Climate Resilience Observatory (CRO), IIIT Lucknow

Lucknow, India

- Conducted EDA on live meteorological data from NASA POWER and ECMWF ERA5 to identify precipitation patterns across 72 Uttar Pradesh districts
- Performed model selection and analytical evaluation for binary rainfall forecasting in collaboration with the Office of the Relief Commissioner, Govt. of Uttar Pradesh

PROJECTS

WealthOS — Personal Financial Intelligence Platform  | LangGraph, FastAPI, Qdrant, DSPy, E2B, Mem0

- Orchestrated **8 agents** via a LangGraph 8-node state machine (**45 tools across 7 MCP servers** — market data, SEC, news, tax, portfolio); `asyncio.gather` parallelism delivers an investment memo in **60–90s** (~2× over sequential) on Groq llama-3.3-70b-versatile
- Built a hybrid **Agentic RAG** over SEC 10-K filings in Qdrant (all-MiniLM dense + BM25 sparse via reciprocal-rank fusion, Cohere rerank, \$0 embedding cost); added **DSPy**-compiled prompts, **E2B**-sandboxed DCF + Monte Carlo, and a 5-dim Financial Health Score — traced in LangSmith + W&B Weave

CalloS — Self-Improving AI Voice Agent Platform  | Google ADK, MCP, A2A, TRL (DPO), DeepEval, GCP

- Built a multi-agent voice platform on **Google ADK** — 7 agents over **6 MCP servers (29 tools)** with A2A endpoints and human-in-the-loop escalation (`asyncio.Future` + `contextvars`); grounded replies via two-stage RAG (MiniLM → cross-encoder rerank) and an LLM-as-judge scorer wired into a DeepEval CI gate
- Implemented a **self-improving fine-tune loop** — filters calls by judge score (≥ 80) into SFT + DPO datasets and runs **DPO alignment** (TRL, Qwen2.5-7B 4-bit QLoRA) on a weekly APScheduler job with a before/after eval harness gating deploy, shipped to GCP Cloud Run via Terraform

RetailGraph — Grocery Intelligence Platform  | Qwen2-VL-7B, QLoRA, LangGraph, GraphRAG, Neo4j, Groq

- Fine-tuned **Qwen2-VL-7B** (QLoRA, 4-bit) on **5,369 verified pairs** (Modal A100), hitting **94.2% exact-match** across 11 fields with 0 parse errors; served via a **6-node LangGraph agent** routing queries to Cypher, GraphRAG (Qdrant + Neo4j), or analytics on FastAPI + Groq Llama-3.3-70B
- Benchmarked **GraphRAG vs vector-only** on 20 queries — vector scored **0/20** on constraints — and scored extraction semantics via **LLM-as-Judge (GPT-4o-mini)**, surfacing 38 hits exact-match missed

SKILLS

Agentic AI: Google ADK, MCP, A2A Protocol, LangGraph, LangChain, Agentic RAG, GraphRAG, DSPy, Mem0

LLM & Fine-tuning: DPO, GRPO, QLoRA, Unsloth, HuggingFace, instructor, LLM-as-Judge, Hybrid Search, Prompt Engineering, Vision-Language Models (Qwen2-VL)

Languages: Python (PyTorch, NumPy, Pandas, Scikit-learn, Matplotlib), C, SQL

Mathematics & Statistics: Linear Algebra, Calculus, Probability Theory, Statistical Inference, Hypothesis Testing

Machine Learning: Transformers, Weak Supervision (Snorkel), Supervised & Unsupervised Learning, CNN, LSTM

Databases & Search: Neo4j (Cypher), Qdrant, pgvector, FAISS, Redis, PostgreSQL

Developer Tools: GCP (Vertex AI, Cloud Run), AWS, Terraform, Docker, FastAPI, Streamlit, LangSmith, W&B, Modal, MLflow

ACHIEVEMENTS

- IIT JAM MS:** Secured All India Rank **219** in IIT JAM 2025 (Mathematical Statistics)
- GATE Statistics:** Secured All India Rank **493** in GATE 2025